

GRLIB VHDL IP Core Library

Configuration and Development Guide

Oct 2025, Version 2025.2

LEON/GRLIB Guide

Table of contents

1	Introduction.....	4
1.1	Overview	4
1.2	Other Resources	4
1.3	Licensing	4
2	System Design Guidelines	5
2.1	Introduction	5
2.2	Minimal System	5
2.2.1	Minimal systems for LEON5 and NOEL-V	5
2.3	Memory Map	6
2.3.1	Overview	6
2.3.2	Typical LEON/GRLIB Memory Map	6
2.3.3	Memory Map in Systems That Need 2 GiB Memory Area	7
2.3.4	LEON5 and NOEL-V memory maps	7
2.3.5	AHB I/O Area and GRLIB Plug&Play Areas	7
2.4	Interrupt Assignments	8
2.4.1	Overview	8
2.4.2	Linux 2.6 and later	8
2.4.3	RTEMS	8
2.4.4	VxWorks	8
2.5	Device Specific Identification	8
3	LEON design information.....	10
3.1	Introduction	10
3.2	General Recommendations.....	10
3.2.1	Data Cache Snooping	10
3.2.2	V7 and FPU	10
3.2.3	MMU and Supervisor Tag bit	10
3.3	LEON Example Configurations	10
3.3.1	Overview	10
3.3.2	Minimal LEON Configuration	11
3.3.3	General Purpose LEON Configuration	12
3.3.4	High Performance LEON Configuration	12
3.3.5	Configuration Settings For Existing LEON Devices	14
3.4	LEON3/4 subsystem (gaisler.subsys.leon_dsu_stat_base)	14
3.5	LEON5 subsystem (gaisler.leon5sys) / NOEL-V subsystem (gaisler.noelvsys)	15
4	Multiple Buses, Clock Domains and Clock Gating	16
4.1	Introduction	16
4.2	Creating Multi-Bus Systems	16
4.2.1	Overview	16
4.2.2	GRLIB Facilities	16
4.2.3	GRLIB AMBA Plug&Play in Multi-Bus Systems	16
4.2.4	Buses in Different Clock Domains	17
4.2.5	Single AHB Bus Example	17
4.2.6	Multi-Bus System Example	17
4.3	LEON3 Double-Clocking	18
4.3.1	Overview	18
4.3.2	LEON3-CLK2X Template Design	18
4.3.3	Clocking	18
4.3.4	Multicycle Paths	19
4.3.5	Dynamic Clock Switching	21
4.3.6	Configuration	21
4.4	Clock gating	21

LEON/GRLIB Guide

	4.4.1	Overview	21
	4.4.2	LEON clock gating	21
5		Debug communication links	23
	5.1	Overview	23
	5.2	Available debug link controllers.....	23
6		Core specific design recommendations	24
	6.1	Overview	24
	6.2	AHB/AHB Bridges (AHB2AHB/AHBBRIDGE/GRIOMMU).....	24
	6.3	SVGA Controller (SVGACTRL)	24
7		GRLIB AMBA Test Framework	25
	7.1	Overview	25
	7.2	AT AHB Master.....	25
	7.2.1	Description	25
	7.2.2	Initialization and Instantiation	25
	7.2.3	Simple Accesses.....	26
	7.3	AT AHB Slave	27
	7.3.1	Description	27
	7.3.2	Initialization and Instantiation	27
	7.3.3	Controlling AT_AHB_SLV	29
	7.4	AT AHB Controller.....	31
	7.4.1	Description.....	31
	7.4.2	Usage.....	31
8		Support.....	32

LEON/GRLIB Guide

1 Introduction

1.1 Overview

The GRLIB IP Library is an integrated set of reusable IP cores, designed for *system-on-chip* (SoC) development. The IP cores are centered around a common on-chip bus, and use a coherent method for simulation and synthesis. The library is vendor independent, with support for different CAD tools and target technologies. A unique plug&play method is used to configure and connect the IP cores without the need to modify any global resources.

The LEON3, LEON4 and LEON5 processors are synthesizable VHDL models of a 32-bit processor compliant with the SPARC V8 architecture. The NOEL-V processor is a synthesizable VHDL model of a 64-bit processor compliant with the RISC-V architecture. The models are highly configurable and particularly suitable for SoC designs. The processor models are distributed as integrated parts of the GRLIB IP Library.

This configuration and development guide is intended to aid designers when developing systems based on LEON/GRLIB. The guide complements the GRLIB IP Library User's Manual and the GRLIB IP Core User's Manual. While the IP Library user's manual is suited for RTL designs and the IP Core user's manual is suited for instantiation and usage of specific cores, this guide aims to help designers make decisions in the specification stage.

1.2 Other Resources

There are several documents that together describe the GRLIB IP Library and Frontgrade Gaisler's IP cores:

- GRLIB IP Library User's Manual (grib.pdf) - Main GRLIB document that describes the library infrastructure, organization, tool support and on-chip bus.
- GRLIB IP Core User's Manual (grip.pdf) - Describes specific IP cores provided with the GRLIB IP library. Also specifies which cores that are included in each type of GRLIB distribution.
- GRLIB-FT User's Manual (grib-ft.pdf) - Describes the FT and FT-FPGA versions of the GRLIB IP library. The document is an addendum to the GRLIB IP Library User's Manual. This document is only available in the FT and FT-FPGA distributions of GRLIB.
- GRLIB FT-FPGA Xilinx Add-on User's Manual (grib-ft-fpga-xilinx.pdf) - Describes functionality of the Virtex5-QV and Xilinx TMRTool add-on package to the FT-FPGA version of the GRLIB IP library. The document should be read as an addendum to the 'GRLIB IP Library User's Manual' and to the GRLIB FT-FPGA User's Manual. This document is only available as part of the add-on package for FT-FPGA.

1.3 Licensing

Some of the cores mentioned in this document (such as LEON4 and the AHB bridges) are only available in the commercial versions of GRLIB.

LEON/GRLIB Guide

2 System Design Guidelines

2.1 Introduction

The design and partitioning of a system strongly depends on the intended use for the system. The sections below make general recommendations based on the components available in GRLIB.

2.2 Minimal System

A minimal LEON/GRLIB system consists of the following IP cores:

TABLE 1. Minimal LEON system

Core	Description
CLKGEN	Clock generator
RSTGEN	Reset generator. Generating a glitch free on-chip system reset signal.
AHBCTRL	AHB arbiter/controller.
APBCTRL	AHB/APB bridge/controller. Must be included in order to interface peripheral cores such as interrupt controller and timer unit.
LEON3/4	LEON3/4 processor
IRQMP	Interrupt controller
GPTIMER	General Purpose Timer Unit
MEMCTRL	Memory controller providing access to (P)ROM and RAM. The GRLIB IP Library contains several memory controllers. It is also possible to include on-chip ROM and RAM by using the AHBROM and AHB RAM IP cores.

In addition to the cores described above it is recommended to include a LEON Debug Support Unit (DSU) and a debug communication link to be able to control the processor and inspect the system via the GRMON Debug Monitor. GRLIB contains several debug communication link (DCL) cores. All DCL cores are controlled over an external link to make accesses on an on-chip AHB bus. Examples of DCL cores are the AHB JTAG, AHB UART and USB DCL cores. See section 5 for more information.

In order for the processor to be able to communicate with the outside world, an 8-bit UART and a General Purpose I/O port is also typically included in a LEON design.

With the above considerations the recommended minimal LEON/GRLIB system also includes the following cores:

TABLE 2. Additional recommended cores for minimal LEON system

Core	Description
DSU3/4	LEON Debug Support Unit
AHB JTAG/ AHB UART/ USB DCL/ GRETH	Debug communication link. AHB JTAG provides an external JTAG link. Other examples include AHB UART (serial UART), USB DCL (USB), GRETH (Ethernet debug communication link is available as part of Ethernet MAC core).
APB UART	8-bit UART
GRGPIO	General Purpose I/O Port

2.2.1 Minimal systems for LEON5 and NOEL-V

LEON5 and NOEL-V are delivered as part of subsystems comprising a minimal system, with exception for the memory controller. Both subsystems includes an AHB system and an APB system, which are exported outside the subsystem, to allow peripherals to be connected to these.

LEON/GRLIB Guide

2.3 Memory Map

2.3.1 Overview

Most LEON systems use a memory map where ROM (boot PROM) is mapped at address 0x00000000 and RAM is mapped at address 0x40000000. Traditionally the AHB/APB bridge has been mapped at 0x80000000 and peripherals such as timer, interrupt controller and UART have been placed at fixed offsets in the APB address space. Table 3 shows the base addresses historically used in LEON systems.

TABLE 3. Peripheral base addresses, legacy systems

Base address	Description
0x80000000	LEON2 memory controller
0x80000100	Generic UART (APBUART)
0x80000200	Multi-processor interrupt controller (IRQMP)
0x80000300	Modular timer unit (GPTIMER)

Some software may not read all peripheral core base addresses from plug&play and instead assume that some peripherals are mapped at these fixed offsets. One of the affected software packages is the BCC 1.0.x toolchain, where the -qambapp switch must be given in order for the produced software to find the UART, timer and interrupt controller in case these peripherals are not mapped at the addresses given in table 3.

The traditional memory map described above does not fit all systems. In particular one or several large memory area (≥ 1 GiB) may be difficult to place as the standard AHB decoder in GRLIB constrains the base address of a memory area based on the memory area size. Other reasons include that the use of AHB-to-AHB bridges that limit how the memory areas can be arranged. As a result of this, there are several LEON/GRLIB designs with different memory maps. In order to ease software development, this document contains some recommendations on how memory maps should be arranged. Section 2.3.2 shows a traditional LEON/GRLIB memory map and section 2.3.3 contains recommendations on how to arrange memory maps that contains large memory areas.

2.3.2 Typical LEON/GRLIB Memory Map

In order to use toolchains and other software distributed by Frontgrade Gaisler, some constraints in the system's memory map should be observed. A typical LEON3 system has the following memory map:

TABLE 4. Typical LEON3 memory map

Base address	Description
0x00000000	PROM
0x40000000	RAM base address. Some systems place SRAM at address 0x40000000 and SDRAM at base address 0x60000000. When SRAM is disabled the memory controller may automatically adjust the SDRAM base address to 0x40000000.
0x80000000	Base address of first AHB/APB bridge connecting interrupt controller, UART(s) and timer unit.
0x90000000	Debug Support Unit register interface
0xFFFF0000	AHB I/O area (if used by any core)
0xFFFFF000	Plug'n'play area (always located within AHB I/O area)

The most important areas in the table above are base addresses for ROM and RAM. The default linker scripts make assumptions on the locations of these areas. Also, software that makes use of the GRLIB AMBA plug'n'play areas often assume the main plug'n'play area to be located at 0xFFFFF000. The

LEON/GRLIB Guide

information in this area is used by software to dynamically find the addresses of all peripherals in the system.

The location of the first AHB/APB bridge (0x80000000 in the table above) is generally of less importance. Some legacy software may assume that the bridge is located at the specified address.

The typical memory map given above constrains the maximum size of a memory area in the design. The GRLIB infrastructure requires that memory areas are binary aligned according to their size. This means that a 2 GiB memory area must start on address 0x00000000 or address 0x80000000. In order to accommodate memory areas of 2 GiB some systems use variations of the memory map as shown in table 5.

2.3.3 Memory Map in Systems That Need 2 GiB Memory Area

TABLE 5. Memory map accomodating 2 GiB main memory area

Base address	Description
0x00000000	RAM
0x80000000	Other large area, for instance PCI bridge mapping PCI memory
0xC0000000	PROM / Memory mapped IO
0xD0000000	AHB/APB bridge
0xE0000000	Debug Support Unit register interface
0xFFFF0000	AHB I/O area (if used by any core)
0xFFFFF000	Plug'n'play area (always located within AHB I/O area)

The memory map in table 5 allows a 2 GiB memory map in the address range 0x00000000 - 0x7FFFFFFF and is supported by the toolchains supplied by Frontgrade Gaisler by giving an extra switch (see the toolchain and OS documentation for details). Note that the default start address for a LEON processor is 0x0. If the memory map above is used, the reset start address should be changed to 0xC0000000.

Existing LEON systems use variations of the above memory map. The main difficulties that can arise from different memory maps is that the RAM and ROM areas may collide in linker scripts and boot loaders. It is therefore recommended that RAM is always mapped at 0x40000000 or 0x00000000 and that ROM (boot PROM area) is mapped at 0x00000000 or 0xC0000000.

Special switches may be required when building the application if RAM is mapped at 0x00000000. See toolchain documentation for details.

2.3.4 LEON5 and NOEL-V memory maps

LEON5 and NOEL-V are delivered as parts of subsystems. The subsystems include an AHB system and an APB system with their default memory maps. For further information about the memory map of LEON5 and NOEL-V systems, refer to the GRLIB IP Core User's Manual (grip.pdf) sections about LEON5SYS and NOELVSYS.

2.3.5 AHB I/O Area and GRLIB Plug&Play Areas

It is recommended that the default addresses are used for AHB I/O areas (determined by generic on AHBCTRL) and GRLIB AMBA plug&play areas (determined via generics on AHBCTRL and APBCTRL). Software scanning routines will assume that one plug&play area is located at 0xFFFFF000.

It is possible to place the AHB I/O area and the AHB plug&play area so that it shadows another AHB area. As an example a PCI core can be mapped at address 0xC0000000 - 0xFFFFFFFF while the plug&play area is still reachable at offset 0xFFFFF000. While such memory maps are perfectly valid and useful for many systems it generally not recommended to let the AHB I/O or plug&play area shadow another area as software drivers may not recognize that some of the memory area assigned to a core is essentially unreachable. When an AHB I/O area or the plug&play area shadows another AHB slave it means that the AHB slave will not be selected when an access is made to the address range occupied by AHB I/O or plug&play.

2.4 Interrupt Assignments

2.4.1 Overview

The LEON processor and interrupt controller provides 15 interrupt lines in the default configuration. Interrupt 15 is non-maskable, which leaves 14 interrupts usable for peripheral cores. The multiprocessor interrupt controllers (IRQMP and IRQ(A)MP cores) can be extended to provide 16 additional interrupts, called extended interrupts.

The GRLIB interrupt infrastructure allows any number of cores to share the same interrupt line. Note, however, that sharing interrupts requires that the software drivers can handle shared interrupts. Also, the time required to serve an interrupt request may be significantly prolonged if software needs to check a large number of registers in order to determine if a peripheral asserted an interrupt.

Some operating systems place additional constraints on interrupt assignments. The subsections below describe the requirements of each OS. The basic rules to follow in order to be able to run the maximum amount of software can be summarized with:

1. If possible, have one dedicated interrupt for each interrupt source (no shared interrupts).
2. Configure the timer unit (GPTIMER) to have dedicated interrupts for each timer
3. Place the timer interrupts within the range 2 - 12
4. Leave interrupts lines 13 - 15 unused

The subsections below dealing with operating systems may become outdated due to changes in the operating systems. If in doubt, please consult the OS documentation or contact Frontgrade Gaisler for the latest information.

2.4.2 Linux 2.6 and later

Interrupt 15 is used for cross-calls. Interrupt 13 is the default selection for inter-processor-interrupts (IPI). The interrupt line to be used for IPI can be selected when building the kernel and cannot be shared with peripherals.

Linux also requires that the first timer on the general purpose timer unit (GPTIMER) has a dedicated interrupt. For SMP operation the second timer must also have a dedicated interrupt line allocated.

2.4.3 RTEMS

RTEMS supports extended interrupts. Interrupt 14 is used for cross-CPU messaging in AMP systems. This interrupt is defined in `leon.h`: `LEON3_MP_IRQ`, cannot be a shared interrupt and must be in the range 1 .. 14.

RTEMS SMP is at the time of writing not finished and requirements are not known.

Timer 0 of GPTIMER 0 is the system clock timer, however RTEMS can be used without a timer. There are two cases depending on which RTEMS distribution that is used:

Classical/official RTEMS BSP: GPTIMER0.timer0 must have separate IRQ and the interrupt must be in the range 1 .. 14.

“Driver manager BSP” (RCC LEON3/4 BSP): Can handle both separate and shared IRQs on GPTIMER, interrupt can be in the range 1 .. 31 (no limitations).

2.4.4 VxWorks

VxWorks makes use of interrupt 14 for inter-processor-interrupts (IPI). This interrupt should not be shared with peripherals.

2.5 Device Specific Identification

GRLIB systems have two identifiers in the system’s plug&play area that can be used to distinguish a particular device: The GRLIB build ID and the GRLIB System Device ID. The GRLIB build ID is set globally for the full library and the device ID is set per design via the `AHBCTRL` VHDL generic `devid` (refer to the `AHBCTRL` section in GRLIB IP Core User’s Manual, `grip.pdf`). This VHDL generic should be set to a unique value for all new designs. The file `lib/grlib/amba/devices.vhd` lists device IDs, under the comment *grlib system device id’s*, used for some existing designs. It is recommended that customer designs use an ID larger than 16#0a00#. Please contact Frontgrade Gaisler support if you wish to have your device ID added to the listing in `devices.vhd`.

LEON/GRLIB Guide

Communication interfaces may have additional vendor and device identifiers. This is, for instance, the case for JTAG, PCI and USB. For the USB debug link it is recommended that users keep the Frontgrade Gaisler IDs so that GRMON may properly detect the debug link. For all other identifiers the implementers of a device should use their own IDs as assigned by the appropriate organisations. Re-use of Frontgrade Gaisler's vendor/manufacture ID may prevent the device from fully functioning together with software and debug tools.

LEON/GRLIB Guide

3 LEON design information

3.1 Introduction

The sections below contain recommendations on how to configure the LEON processors depending on system requirements.

3.2 General Recommendations

3.2.1 Data Cache Snooping

To keep the data cache synchronized with external memory, cache snooping can be enabled. When enabled, the data cache monitors write accesses on the AHB bus to cacheable locations. If another AHB master writes to a cacheable location that is currently cached in the data cache, the corresponding cache line is marked as invalid.

Data cache snooping is of high importance for SMP systems and, in general, both simplifies and increases performance in systems with multiple masters. Note that the processor(s) snoop on the bus to which they are directly connected. In a system with multiple AHB buses, snooping will only work on the bus to which the processors are connected. Snooping will not provide cache coherency if, for instance, there are masters connected between a Level-2 cache and memory, while the processors are located in front of the Level-2 cache.

Snooping is also required to prevent aliasing effects in systems that use a memory management unit (MMU) and a data cache that is either larger than the MMU page size or has more than one way. If the processor(s) is implemented with a MMU, then separate snoop/physical tags must be enabled. A multi-way cache allows virtual addresses mapped to the same physical address to be cached in each cache-way. A write operation will only update the copy in one of the cache-ways, leading to data coherency issues. If snooping with separate physical tags is enabled then the aliased addresses will be invalidated by the processor write operation.

In case the data cache way size is larger than the MMU page size then multiple MMU pages may be cached in the same cache-way, again leading to the cache containing multiple virtual locations that map to the same physical address. In this case the LEON snooping implementation will not resolve the situation and the processor needs to be implemented with a MMU page sizes that matches the cache way size. Note that operating system support for MMU page sizes larger than 4 KiB is limited, consult the documentation for the operating system.

3.2.2 V7 and FPU

When the LEON is implemented with an FPU it should also include hardware support for multiply and divide (SPARC V8 MUL/DIV selected with the LEON VHDL generic `v8`). Otherwise a SPARC V7 processor with FPU will be obtained and this configuration may not be supported by prebuilt packages and toolchains.

3.2.3 MMU and Supervisor Tag bit

When the LEON is implemented with an MMU it is recommended to include the supervisor access bit in the L1 cache tag. Otherwise there is a risk of information leaking to user mode from kernel mode due to kernel data being present and accessible from user space in the L1 cache. The extra tag information is included by setting the `mmuen` VHDL generic to 2. The extra tag bit does not provide any extra functionality for systems that only use supervisor mode and use the MMU as an extra safety net, in these cases the bit can be disabled to reduce the width of L1 tag RAMs with one bit.

3.3 LEON Example Configurations

3.3.1 Overview

The subsections below show three different example configurations for LEON processors; a minimal configuration used to target low area and high frequency, a typical configuration with all features enabled, and a high-performance configuration where the requirements on processing performance outweigh area and power considerations.

Each section contains a table with recommended values for some of the LEON processor VHDL generics. If you are using the **xconfig** GUI to configure the processor then please note that the VHDL

generic names do not directly correspond to the configuration options in the GUI. The descriptions of the configuration settings should provide enough information to do appropriate configuration selection also via **xconfig**. The xconfig tool also has support to initialize the processor configuration with values from the three example configurations described in the sections below. See the configuration help text in **xconfig** for the option *Force values from example configuration* in the *Processor* sub menu for additional information.

Also note that all listed configuration options do not apply to all LEON processors. For instance, the LEON3 processor has a VHDL generic called bp that controls the inclusion of branch predication, while the LEON4 processor is always implemented with support for branch prediction.

3.3.2 Minimal LEON Configuration

This LEON configuration is aimed at resource constrained systems where the area requirements of the processor core needs to be minimized. Note that using an area minimized configuration may not necessarily reduce the system's performance since it may be possible to achieve a higher operating frequency by reducing the amount of logic in the processor core.

Table 6 below shows recommended values for some of the LEON processor VHDL generics to attain a minimal configuration in terms of area.

TABLE 6. Minimal LEON processor configuration

VHDL generic	Recommended value	Description
dsu	0	Some area can be saved by removing the Debug Support Unit (DSU). However, this unit can prove to be invaluable at least during the software development phase.
fpu	0	Disable floating-point unit
v8	0	Do not include support for SPARC V8 MUL/DIV instructions
mac	0	Do not include support for SPARC V8e SMAC/UMAC
nwp	0	Disable hardware watchpoints
icen / dcen	1	Include processor caches
isets / dsets	1	Direct mapped instruction and data cache
irepl / drepl	2	Random replacement policy for both instruction and data cache (setting is unused for direct-mapped cache)
isetsize / dsetsize	-	The size of the caches does not significantly affect the required logic. Choose cache size according to application requirements and amount of RAM available on target device.
dnsnoop	0	Disable data cache snooping (see section 3.2.1)
mmuen	0	Disable memory management unit (MMU). Note: May be required depending on software applications.
lddel	1	1-cycle load delay
tbuf	0	Disable instruction trace buffer (NOTE: Including the instruction trace buffer may be of high value during software development and debug).
pwd	1	Power-down implementation. Choose 2 if frequency target is not met.
smp	0	Disable SMP support. If the processor core should be used in an SMP configuration then see the GRIP documentation on how to set the SMP generic. If SMP is enabled then the <i>dsnoop</i> VHDL generic should also be set accordingly.
bp	0	Disable branch prediction

LEON/GRLIB Guide

3.3.3 General Purpose LEON Configuration

This LEON configuration is aimed for general purpose processing balancing performance against area and power requirements.

TABLE 7. General purpose LEON processor configuration

VHDL generic	Recommended value	Description
dsu	1	Include support for the LEON Debug Support Unit (DSU)
fpu	-	Include floating-point unit based on application requirements. A floating-point unit is highly recommended for most systems. LEON processors can primarily interface the GRFPU or GRFPU-lite floating point unit. The GRFPU is a high-performance pipelined FPU with high area requirements. GRFPU-lite provides a balanced option with high acceleration of floating-point computations combined with lower area requirements compared to GRFPU.
v8	2	Include support for SPARC V8 MUL/DIV instructions using a 5-cycle multiplier. Note that if the target technology has multiplier blocks a single-cycle multiplier (v8 generic set to 1) may provide lower area and higher performance.
mac	0	Do not include support for SPARC V8e SMAC/UMAC instructions.
nwp	2	Include two hardware watchpoints
icen / dcen	1	Include processor caches.
isets / dsets	2	Implement instruction and data caches with two ways
irepl / drepl	2	Random replacement policy for both instruction and data cache, or possibly LRU replacement (irepl/drepl set to 0).
isetsize / dsetsize	-	The size of the caches does not significantly affect the required logic. Choose cache size according to application requirements and amount of RAM available on target device.
dnsoop	6	Enable snooping with extra physical tags (see section 3.2.1)
mmuen	2	Enable memory management unit (MMU)
itlbnun / dtlbnun	8	Use eight entries each for the instruction and data MMU translation look-a-side buffers
tlb_type	2	Use separate translation look-a-side buffers (TLB) with fast write for data and instruction.
tlb_rep	0	Use LRU TLB replacement
lddel	1	Use 1-cycle load delay
tbuf	4	Use 4 KiB instruction trace buffer.
pwd	2	Timing efficient power-down implementation.
smp	0	Disable SMP support. If the processor core should be used in an SMP configuration then see the GRIP documentation on how to set the SMP generic.
bp	1	Enable branch prediction

3.3.4 High Performance LEON Configuration

This LEON configuration is aimed at high performance processing where the needs for computational speed outweighs area and power requirements.

LEON/GRLIB Guide

In order to reduce the effects of memory latency, a Level-2 cache is recommended for high-performance systems. This is of particular interest in multiprocessor systems.

TABLE 8. High-performance LEON processor configuration

VHDL generic	Recommended value	Description
dsu	1	Include support for the LEON Debug Support Unit (DSU)
fpu	1 - 7	Use GRFPU floating-point unit. Select (FP) multiplier depending on target technology. For FPGA this would typically be inferred (1) or technology specific (4). For ASIC DesignWare multiplier (2) or Module Generator (3).
v8	16#32#	Include support for SPARC V8 MUL/DIV instructions using a 32x32 pipelined multiplier. Note that if the target technology has multiplier blocks a single-cycle multiplier (v8 generic set to 1) may provide lower area and higher performance.
mac	0	Do not include support for SPARC V8e SMAC/UMAC instructions
nwp	4	Include support for four hardware watchpoints
icen / dcen	1	Include processor caches.
isets / dsets	2	Implement instruction and data caches with two ways
irepl / drepl	0	Least-Recently-Used replacement policy for instruction and data caches.
isetsize / dsetsize	-	The size of the caches does not significantly affect the required logic. Choose cache size according to application requirements and amount of RAM available on target device.
dnsoop	6	Enable snooping with extra physical tags (see section 3.2.1)
mmuen	2	Enable memory management unit (MMU)
itlbnum / dtlbnum	16	Use sixteen entries each for the instruction and data MMU translation look-a-side buffers
tlb_type	2	Use separate translation look-a-side buffers (TLB) with fast write for data and instruction.
tlb_rep	0	Use LRU TLB replacement
lddel	1	Use 1-cycle load delay
tbuf	4	Use 4 KiB instruction trace buffer.
pwd	2	Timing efficient power-down implementation.
smp	> 0	Enable SMP support. If the processor core should be used in an SMP configuration then see the GRIP documentation on how to set the SMP generic. Note that several processor entities must be instantiated. This configuration option only enables support for SMP, it does not instantiate several processor cores.
bp	1	Enable branch prediction

LEON/GRLIB Guide

3.3.5 Configuration Settings For Existing LEON Devices

The table below shows configurations for existing Frontgrade LEON devices. Please refer to the previous subsections under section 3.3 for comments and descriptions of the different values.

TABLE 9. LEON processor configurations

VHDL generic	UT699 Value	UT700 Value	GR712RC value	GR740 Value	LEON3-RTAX example value
dsu	1	1	1	1	1
fpu	2	2	2	2	0
v8	2	16#32#+4	2	16#32#	0
mac	0	0	0	0	0
nwp	4	4	2	4	2
icen	1	1	1	1	1
isets	2	4	4	4	1
isetsize	4	4	4	4	8
irepl	0	0	0	0	0
ilinesize	8	8	8	8	8
dcen	1	1	1	1	1
dsets	2	4	4	4	1
dsetsize	4	4	4	4	4
drepl	0	0	0	0	0
dlinesize	4	4	4	8	4
dnsoop	6	6	6	6	0
mmuen	1	1	1	1	0
itlbnm / dtlbnm	16 / 16	16 / 16	16 / 16	16 / 16	- / -
tlb_type	0	2	2	2	0
tlb_rep	0	0	0	0	0
lddel	2	1	1	1	1
tbuf	2	4	4	8	2
pwd	2	2	2	2	2
svt	1	1	1	1	1
smp	0	0	1	1	0
bp	N/A (0)	1	1	N/A	0
npasi	N/A (0)	N/A (0)	N/A (0)	1	N/A (0)
pwrpsr	N/A (0)	N/A (0)	N/A (0)	1	N/A (0)
LEON version used	LEON3FT v1	LEON3FT v2	LEON3FT v1 with BP	LEON4v0	LEON3FTv1 to LEON3v3

3.4 LEON3/4 subsystem (gaisler.subsys.leon_dsu_stat_base)

GRLIB contains a subsystem component that can be used to instantiate the LEON processor, debug support unit and a statistics unit (performance counters). The subsystem is available in *lib/gaisler/subsys/* and also has a corresponding **xconfig** script. Please refer to GRLIB IP Core User's Manual (grip.pdf) for documentation of LEON_DSU_STAT_BASE.

3.5 LEON5 subsystem (gaisler.leon5sys) / NOEL-V subsystem (gaisler.noelvsys)

Both LEON5 and NOEL-V are delivered as part of a subsystem, designed to provide a necessary minimal LEON5/NOEL-V system. Both subsystems are configurable to allow tailoring of the functionality to meet performance and resource utilization goals. The subsystem for LEON5 is available in *lib/gaisler/leon5v0* and the subsystem for NOEL-V is available in *lib/gaisler/noelv/subsys*. Please refer to GRLIB IP Core User's Manual (grip.pdf) for documentation about LEON5SYS and NOELVSYS.

4 Multiple Buses, Clock Domains and Clock Gating

4.1 Introduction

This section describes some techniques that can be used with GRLIB to create more complex system architectures with multiple buses and/or clock domains.

Peripheral IP cores that need to work at a separate clock domain usually have their own clocking and synchronization built in. This is not explained here, see the core-specific documentation.

4.2 Creating Multi-Bus Systems

4.2.1 Overview

The on-chip bus may become a bottle neck in systems where the processors and peripherals all share the same bus. The fact that all IP cores are connected together may also introduce high loads in the system, which can lead to timing issues at implementation. These issues can be solved by partitioning the system into several AHB buses.

4.2.2 GRLIB Facilities

In order to partition the system into multiple buses, the general-purpose AHB bridge IP cores AHB-BRIDGE (uni-directional) and AHB2AHB (bi-directional) are included in GRLIB. There are also special-purpose cores, such as the IOMMU and L2-cache, that have bridge functionality built into them.

4.2.3 GRLIB AMBA Plug&Play in Multi-Bus Systems

Software and debug monitors such as GRMON can detect all IP cores connected to the on-chip bus(es) by scanning the plug&play configuration area. The format and function of this area is described in the GRLIB User's Manual and in the GRLIB IP Core User's Manual documentation for the AHB controller (AHBCTRL) and AHB/APB bridge (APBCTRL).

In multi-bus systems, each bus will have its own AMBA plug&play configuration area and software must be able to access all plug&play areas. In order for software to be able to discover all peripherals in a system. The same applies for the GRMON debug monitor, to discover all peripherals the debug communication link master interface must be connected to a bus from where it can access all plug&play areas (as well as memory where peripheral registers are mapped).

The plug&play scanning routines discover the presence of multiple AHB buses when it discovers the slave interface a core such as the Level-2 cache or AHB/AHB bridge (AHB2AHB, AHB-BRIDGE). Upon discovery of a bridge the routine will typically look in the user defined register of the bridge's plug&play information to get the base address of the AHB I/O and plug&play area of the second bus. Exactly how the base address of the plug&play information is communicated to the scanning routine is specific for each core. The Level-2 cache and AHB/AHB bridges store this address in user defined register 1 of the core's AHB slave interface plug&play information. A value of zero in this register signals to software that plug&play scanning should not be done for the second bus behind the bridge.

When software discovers a bridge to a new bus, scanning should commence using the new plug&play area address (depth-first scanning) and once the new plug&play area has been handled scanning should continue on the current bus.

Note that for plug&play scanning to work, all plug&play areas must be accessible from the AHB master that performs the scan. This means that any bridge between AHB buses must have a window that allows the plug&play area on the other side of the bridge to be accessed. System software and debug tools by default start scanning for a plug&play area at the top of AMBA memory space. It is important that the plug&play area located in this address has pointers so that all other plug&play areas in the system can be discovered. For instance, the default plug&play area address should **not** be occupied by the plug&play area of a bus that is **only** connected to the rest of the system via the AHB master interface side of a Level-2 cache or uni-directional bridge. This is because the extra information at the AHB master interface does not contain the base address for the plug&play area of the bus on the AHB slave interface side of the bridge. As a result of this, plug&play scanning routines will only find one bus in the system.

LEON/GRLIB Guide

4.2.4 Buses in Different Clock Domains


In order to work around timing issues, or to reduce power consumption, it can make sense to partition the design also into several clock domains. The AHB/AHB bridges (AHB2AHB, AHBBRIDGE and GRIOMMU) allows connecting buses with differing operating frequencies together.

The bus clocks on each side of the bridge need to have a frequency ratio relationship and fixed phase relation. This avoids the need to resynchronize signals on chip which would cause a performance penalty.

If you want to run everything except the processor at half speed, a more efficient solution than using bridges is to use the LEON double clocking support explained in section 4.3.


4.2.5 Single AHB Bus Example

A typical LEON/GRLIB design is shown in the figure below. The design is centered around one AMBA AHB bus and also has a AMBA APB bus that connects some of the peripheral cores via an AHB/APB bridge.



Building the system around one AHB bus has advantages in that it simplifies system design.

4.2.6 Multi-Bus System Example



One example (shown above) of when a multi-bus system resolves bus contention is when a SVGA controller (SVGACTRL core) is used. The SVGA controller continuously reads a frame buffer located in external memory. This constant data fetching can consume a significant amount of the available bus bandwidth, particularly in systems with relatively low system frequencies. The impact of the SVGA controller bus traffic can be removed by placing the SVGA controller and a dedicated memory controller on a separate bus. The processor can still access the frame buffer through a uni-directional bridge.

4.3 LEON3 Double-Clocking

4.3.1 Overview

To avoid critical timing paths in large AHB systems, it is possible to clock the LEON3 processor core at an integer multiple of the AHB clock. This will allow the processor to reach higher performance while executing out of the caches. The performance will be higher while executing out of the caches since the processor core will be running at a higher frequency. On a cache miss the processor will need to make a bus access and timing of this bus access will be made according to the lower bus frequency. This chapter will describe how to implement a LEON3 double-clocked system using the LEON3-CLK2X template design as an example.

The LEON3 CPU core can be clocked at a multiple of the clock speed of the AMBA AHB bus. When clocked at double AHB clock frequency, all CPU core parts including integer unit and caches will operate at double AHB clock frequency while the AHB bus access is performed at the slower AHB clock frequency. The two clocks have to be synchronous and multicycle paths between the two clock domains have to be defined at synthesis tool level. Separate components (leon3s2x, leon3x, leon3ft2x) are provided for the double clocked core. Double clocked versions of DSU (dsu3_2x) and MP interrupt controller (irqmp2x) are used in a double clocked LEON3 system. An AHB clock qualifier signal (*clken* input) is used to identify end of AHB cycle. The AHB qualifier signal is generated in CPU clock domain and is high during the last CPU clock cycle under AHB clock low-phase.

4.3.2 LEON3-CLK2X Template Design

The LEON3-CLK2X design is a multi frequency design based on double-clocked LEON3 CPU core. The LEON3 CPU core and DSU run at multiple AHB frequency internally, while the AHB bus and other AHB components are clocked by the slower AHB clock. Double clocked version of the interrupt controller is used, synchronizing interrupt level signals between the CPU and the interrupt controller.

The design can be configured to support different ratios between CPU and AHB clock such as 2x, 3x or 4x. If dynamic clock switching is enabled, a glitch-free clock multiplexer selecting between the fast CPU clock and the slower AHB clock is used to dynamically change frequency of the CPU core (by writing to an APB register).

4.3.3 Clocking

The design uses two synchronous clocks, AHB clock and CPU clock. For Xilinx and Altera technologies the clocks are provided by the *clkgen* module, for ASIC technologies a custom clock generation circuit providing two synchronous clocks with low skew has to be provided.

An AHB clock qualifier signal, identifying end of an AHB clock cycle is necessary for correct operation of the double-clocked cores. The AHB clock qualifier signal (HCLKEN), indicating end of an AHB clock cycle, is provided by the *qmod* module. The signal is generated in CPU clock domain and is active during the last CPU clock cycle during low-phase of the AHB clock. Figure 1 shows timing for CPU and AHB clock signals (CPUCLK, HCLK) and AHB clock qualifier signal (HCLKEN) for clock ratios 2x and 3x.




Figure 1. Timing diagram for CPUCLK, HCLK and HCLKEN

4.3.4 Multicycle Paths

Paths going through both CPU and AHB clock domains have propagation time of one AHB clock cycle, and should be marked as multicycle paths with following exceptions:

Start point	Through	End point	Propagation time
leon3s2x core			
CPUCLK	ahbi	CPUCLK	N CPUCLK
CPUCLK	ahbsi	CPUCLK	N CPUCLK
CPUCLK	ahbso	CPUCLK	N CPUCLK
HCLK	irqi	CPUCLK	1 CPUCLK
CPUCLK	irqo	HCLK	1 CPUCLK
CPUCLK		u0_0/p0/c0/sync0/r[*] (register)	1 CPUCLK
dsu3_2x core			
CPUCLK	ahbmi	CPUCLK	N CPUCLK
CPUCLK	ahbsi	CPUCLK	N CPUCLK
	dsui	CPUCLK	1 CPUCLK
r[*] (register)		rh[*] (register)	1 CPUCLK
irqmp2x core			
r2[*] (register)		r[*] (register)	1 CPUCLK

* N is ratio between CPU and AHB clock frequency (2, 3, ...)

Sample DC script defining multicycle paths and exceptions is provided in the design directory (*dblclk.dc*).

Figure 2 shows synchronization of AHB signals starting in HCLK clock domain and ending in CPUCLK domain (inside the double clocked cores LEON3S2X and DSU3_2X). These AHB signals are captured by registers in CPUCLK domain at the end of AHB clock cycle, allowing propagation time of 2 or more CPUCLK cycles (one HCLK cycle). The end of the AHB clock cycle is indicated by the AHB clock qualifier signal HCLKEN. One of the inputs of the AND gate in figure below is connected to the clock qualifier signal HCLKEN ensuring that the value of the signal AHBI is latched into R2 at the end of AHB cycle (HCLKEN = '1'). The value of signal AHBI is not valid in the CPUCLK clock domain if the qualifier signal HCLKEN is low. In this case, the AND gate will be closed and the value of the signal AHBI will not propagate to register R2.




Figure 2. Synchronization between HCLK and CPUCLK clock domains

Synchronization of AHB signals going from the double clocked cores to the AHB clock domain is shown in figure 3. The AND gate is open when CPU (or DSU) performs an AHB access (AHBEN = '1'). When the AND gate is open, the signal AHBO will be stable during the whole AHB cycle and its value propagates to the HCLK clock domain (AHB bus). When CPU does not perform AHB access (AHBEN = '0') the AND gate is closed (AHBEN = '0') disabling propagation of signal AHBO to the HCLK clock domain.




Figure 3. Synchronization between CPUCLK and HCLK clock domains

The AND gates in figures 2 and 3 are 2-input clock AND gates. Synthesis tool should not optimize these AND gates. Sample DC-script puts 'don't-touch' attribute on these cells to prevent optimization.

The multicycle constraints for the GRLIB double clocked cores are typically defined by start clock domain, intermediate points and end clock domain. Although FPGA synthesis tools provide support for multicycle paths, they do not provide or have limited support for this type of multicycle con-

LEON/GRLIB Guide

straints (start clock domain, intermediate points, end clock domain). This limitation results in over-constrained FPGA designs (multicycle paths become single cycle) which are fully functional and suitable for FPGA prototyping.

4.3.5 Dynamic Clock Switching

An optional clock multiplexer switching between the CPU and AHB clocks and providing clock for double-clocked cores can be enabled. The clock multiplexer is used to dynamically change frequency of the CPU core, e.g. CPU can run at lower AHB frequency during periods with low CPU load and at twice the AHB frequency during periods with high CPU load.

The clock switching is controlled by writing to the *qmod* modules APB register (default address 0x80000400), bit 0: writing '1' will switch to the CPU clock and writing '0' will switch to the AHB clock.

The clock multiplexer is glitch-free, during clock switching the deselected clock is turned-off (gated) before the selected clock is enabled and selected.

Dynamic clock switching is available for Xilinx and generic technologies.

4.3.6 Configuration

xconfig

Clock ratios 2x, 3x and 4x between CPU and AHB clock are supported. Clock ratio 2x is supported for all technologies, ratios 3x and 4x are supported for ASIC technologies. Dynamic clock switching is available for Xilinx and ASIC technologies.

leon3s2x

Double-clocked LEON3 core is configured similarly to standard LEON3 core (leon3s) through VHDL generics. An additional VHDL generic *clk2x* is set to $((\text{clock ratio} - 1) + (8 * \text{dyn}))$ where *dyn* is 1 if dynamic clock switching is enabled and 0 if disabled.

qmod

Local *qmod* module generates AHB clock qualifier signal and optionally controls dynamic clock switching. The module is configured through VHDL - generics defining clock ratio (*clkfact*), dynamic clock switching (*dynfreq*) and address mapping of modules APB register (*pindex*, *paddr*, *pmask*).

irqmp_2x

VHDL generic *clkfact* should be set to clock ratio between CPU and AHB clocks.

4.4 Clock gating

4.4.1 Overview

GRLIB contains support for using clock gating for both the processors and peripheral IP cores. The GRCLKGATE unit described in the GRLIB IP Core User's Manual can be used both to gate peripherals and to provide automatic processor (and floating-point unit) clock gating.

4.4.2 LEON clock gating

To further reduce the power consumption of the processor, the clock can be gated-off when the processor has entered power-down state. Since the cache controllers and MMU operate in parallel with the processor, the clock cannot be gated immediately when the processor has entered the power-down state. Instead, a power-down signal (DBG0.idle) is generated when all outstanding AHB accesses have been completed and it is safe to gate the clock. This signal should be clocked through a positive-edge flip-flop followed by a negative-edge flip-flop to guarantee that the clock is gated off during the clock-low phase. To ensure proper start-up state, the clock should not be gated during reset and at least 3 clocks after that reset has been de-asserted.




Figure 4. Examples of LEON clock gating

The processor should exit the power-down state when an interrupt become pending. The signal `DBG0.ipend` will then go high when this happen, and should be used to re-enable the clock.

When the debug support unit (DSU3 or DSU4) is used, the `DSU0.pwd` signal should be used instead of `DBG0.idle`. This will ensure that the clock also is re-enabled when the processor is switched from power-down to debug state by the DSU. The `DSU0.pwd` is a vector with one power-down signal per CPU (for SMP systems). `DSU0.pwd` takes `DBG0.ipend` into account, and no further gating or latching needs to be done of this signal. If cache snooping has been enabled, the continuous clock will ensure that the snooping logic is activated when necessary and will keep the data cache synchronized even when the processor clock is gated-off. In a multi-processor system, all processor except node 0 will enter power-down after reset and will allow immediate clock-gating without additional software support.

Clock-tree routing must ensure that the continuous clock (CLK) and the gated clock (GCLK) are phase-aligned. The template design *leon3-clock-gate* shows an example of a clock-gated system. Please refer to the LEON signal descriptions in the GRLIB IP Core User's Manual document for documentation on which processor clock inputs that are allowed to be gated-off. Please also see the documentation for the `GRCLKGATE` and `GRCLKGATE2` IP cores in the same document.

LEON/GRLIB Guide

5 Debug communication links

5.1 Overview

GRLIB contains several debug communication link (DCL) controller cores. All DCL cores are controlled over an external link to make accesses on an on-chip AHB bus. These communication links can be used by an external debug monitor to perform debugging operations on the system or by other external devices that need direct memory access to the design.

5.2 Available debug link controllers

A debug communication link controller is an IP core that has that supports communication over an external interface without on-chip software involvement. The IP core decodes incoming traffic and translates the traffic to operations on the AMBA bus. The table below lists IP cores that can act as debug communication link controllers.

TABLE 10. Debug Communication Link controllers

Interface	IP core	AMBA access size supported	Notes
Serial UART	AHBUART	Word	Supported by GRMON
JTAG	AHBJTAG	Byte, Half-word, Word	Supported by GRMON
Ethernet	GRETH / GRETH_GBIT	Word	DCL functionality is optional to include in Ethernet controllers. Supported by GRMON.
PCI	GRPCI / GRPCI2	Byte, Half-word, Word	GRMON can make use of PCI target to access system.
SpaceWire RMAP	GRSPW GRSPW2 / GRSPWROUTER	Read: Byte, Half-word, Word Write: Word	RMAP hardware handler is optional to include in SpaceWire controllers. GRMON can connect via GRESB Ethernet-to-SpaceWire bridge. The controllers translate sub-word read accesses to 32-bit read operations.
USB	GRUSB_DCL	Word	Supported by GRMON
I2C	I2C2AHB	Byte, Half-word, Word	Not supported by GRMON
SPI	SPI2AHB	Byte, Half-word, Word	Not supported by GRMON

6 Core specific design recommendations

6.1 Overview

The subsections below contain system design recommendations when using specific GRLIB cores.

6.2 AHB/AHB Bridges (AHB2AHB/AHB2BRIDGE/GRIOMMU)

The AHB/AHB bridges can be of high value when partitioning the system into several clock domains or when there is a need to separate bus traffic. The use of a bridge will result in increased latencies when accesses need to traverse over the bridge.

For bi-directional bridge configurations the designer needs to be aware that collisions (attempts to traverse the bridge both ways simultaneously) will mean that the access on the slave bridge will be aborted and then re-attempted. This situation can potentially lead to starvation and deadlocks.

When instantiating the bridge with a prefetch buffer the buffer should be scaled so that it does not prefetch unnecessarily large amounts of data. If the master(s) traversing the bridge have a maximum burst length of eight words, then the bridge's prefetch buffer should not be larger than eight words.

6.3 SVGA Controller (SVGACTRL)

The SVGA controller can consume a significant amount of the available bus bandwidth. Even if calculations show that there is plenty of bandwidth available, the inclusion of SVGACTRL may add bus access latencies that significantly impact computational performance. For design that include a SVGA controller it is recommended to place the SVGA controller on a separate bus with a dedicated frame buffer memory.

7 GRLIB AMBA Test Framework

7.1 Overview

GRLIB has a number of packages that can aid in verification of AMBA cores. New developments should use the GRLIB AMBA Test Framework (ATF). The test framework consists of an AHB master core, an AHB slave core and an AHB arbiter/controller core. The AHB master and slave cores have debug interfaces that allow them to be controlled using external stimuli.

The sections below give an overview of the components in the framework. The test framework is not distributed as a product and there is no complete user's manual. The test master and slave is controlled by procedure calls that are documented in their respective VHDL packages (described below).

ATF files are located in the directory `<grlib root>/lib/grlib/atf/`. All GRLIB distributions do not include ATF. If the `atf` directory is missing from your GRLIB tree, then your version of GRLIB does not contain the components described in this section.

NOTE: The GRLIB AMBA test framework is NOT included in the free GRLIB-GPL.

7.2 AT AHB Master

7.2.1 Description

The AT AHB Master (AT_AHB_MST) is a non-synthesizable AHB master core with a debug interface so that the master can be controlled via function calls.

7.2.2 Initialization and Instantiation

The component for the master is defined in the package `grlib.at_pkg` and the procedure calls to control the master is available in the package `grlib.at_ahb_mst_pkg`. In order to instantiate the master, the following libraries should be included:

```
library ieee;
use ieee.std_logic_1164.all;
library grlib;
use grlib.amba.all;
use grlib.at_pkg.all;
use grlib.at_ahb_mst_pkg.all;
use grlib.testlib.all;
```

The component for AT_AHB_MST has the following interface:

```
component at_ahb_mst is
  generic(
    hindex:      in    Integer := 0;
    vendorid:    in    Integer := 0;
    deviceid:    in    Integer := 0;
    version:     in    Integer := 0;
    grlibdatamux: in    integer := 1);
  port(
    -- AMBA AHB system signals
    hclk:        in    std_ulogic;
    hresetn:     in    std_ulogic;
    --AHB Interface
    ahbi:        in    ahb_mst_in_type;
    ahbo:        out   ahb_mst_out_type;
    --Operation Scheduling Interface
    atmi:        in    at_ahb_mst_in_type;
    atmo:        out   at_ahb_mst_out_type
  );
end component;
```

The only VHDL generics that require proper assignment are `hindex` and `grlibdatamux`. The `hindex` generic must match the bus index in the same way as for other GRLIB AHB masters. The `grlibdatamux` generic decides if the core should use AMBA compliant data multiplexing (`grlibdatamux => 0`) or the simplified data multiplexing scheme (`grlibdatamux => 1`) commonly used in GRLIB (see the GRLIB IP Library User's Manual, *grlib.pdf*, for details). For use in a normal GRLIB system the default value is recommended. An example instantiation of AT_AHB_MST can be found in *verification/at/at_tb.vhd*. At the top of the file the libraries mentioned above are included. The test bench instantiates several AMBA masters, the signals used to control the debug interfaces are created as:

```
signal atmi : at_ahb_mst_in_vector(0 to 2);
signal atmo : at_ahb_mst_out_vector(0 to 2);
```

LEON/GRLIB Guide

The masters are then instantiated using a generate loop:

```
-- Masters
mstrs01 : for i in 0 to 2 generate
  amst : at_ahb_mst
    generic map(
      hindex      => FIRST_MASTER_INDEX+i,
      vendorid    => 0,
      deviceid    => 0,
      version     => 0)
    port map(
      -- AMBA AHB system signals
      hclk        => clk,
      hresetn     => rstn,

      -- Direct Memory Access Interface
      atmi        => atmi(i),
      atmo        => atmo(i),

      -- AMBA AHB Master Interface
      ahbi        => ahbmi,
      ahbo        => ahbmo(FIRST_MASTER_INDEX+i));
end generate;
```

The masters are controlled by calls from the test bench process. Before use, each master debug interface must be initialized. In *verification/at/at_tb.vhd* this is done by calls to *at_init(..)*:

```
testbench: process

----- variable definitions removed -----

begin -- process testbench
-----
-- Testbench initialization
-----
Print("-----");
Print("AMBA Test Framework test bench");
Print("-----");
for i in atmi'range loop
  at_init(i, atmi);
end loop;
wait until rstn = '1'
```

7.2.3 Simple Accesses

After initialization has been performed, as described in the previous section, the procedures defined in *grrlib.at_ahb_mst_pkg* (*lib/grrlib/atf/at_ahb_mst_pkg.vhd*) can be used to command the master to perform accesses. The procedures are either read or write procedures. A read or write procedure can be either blocking (call will not return before the access is completed) or non-blocking (call will return immediately and another call must be made at a later time in order to complete the command on the debug interface). All non-blocking procedures have names ending with *_nb*, the procedures used to complete a non-blocking call have names that end with *_nb_fin*.

Procedures that make single accesses are named in the following format: *at_read <size>(..)* or *at_write <size>(..)*. Where *<size>* can be 8, 16, 32, 64, 128 or 256. The non-blocking pairs are named *at_read <size>_nb(..)* / *at_read <size>_nb_fin(..)* and *at_write_nb(..)* / *at_write_nb_fin(..)*. There are also procedures that make burst accesses. These have the word *burst* in their name, for instance *at_write_burst_32(..)*. The procedure names are overloaded and there can be several variants of a procedure, with a different number of parameters.

The simplest way to perform a single access, in this case a write, is to use a call like:

```
at_write_32(
  address => X"h40000000",
  data => X"01234567",
  atmi => atmi(0),
  atmo => atmo(0));
```

The non-blocking variant is (here we assume that we have defined the variable *id* as an integer and the variable *ready* as a boolean):

```
at_write_32_nb(
  address => X"h40000000",
  data => X"01234567",
  waitcycles => 0,
  lock => false,
  hprot => "0011",
  back2back => false,
  screenoutput => false,
```

LEON/GRLIB Guide

```

    id => id,
    atmi => atmi(0),
    atmo => atmo(0));

-- Here other tasks can be performed

at_write_32_nb_fin(
    id => id,
    wait_for_op => true,
    screenoutput => false,
    ready => ready,
    atmi => atmi(0),
    atmo => atmo(0));

```

The first call initiates a write access to address 0x40000000 with data 0x01234567. The access should start immediately, not assert HLOCK and use the specified HPROT (0b0011). The first call will assign an access identifier to the variable *id*. This identifier is used by AT_AHB_MST to keep track of the access. The same access identifier must then be used in the call to *at_write_32_nb_fin(..)*. The core will try to perform the write access even if the call to *at_write_32_nb_fin(..)* never takes place. However, if *at_write_32_nb_fin(..)* is never called, the core will keep a record of the completed access in its internal data structures forever.

A call to *at_<operation>_<size>_nb_fin(..)* procedure will block if the *wait_for_op* parameter is set to *true*. If *wait_for_op* is set to *false*, the call will return immediately and the *ready* variable must be checked to see if AT_AHB_MST completed the access.

The description given for write operations above also applies to read operations. Note that for non-blocking reads (*at_read_<size>_nb(..)* / *at_read_<size>_nb_fin(..)*), the data will be returned when *at_read_<size>_nb_fin(..)* is called. The first call only tells the master to initiate an access, the *at_read_<size>_nb_fin(..)* call will tell you when, and if, the access has completed and the master will have data available.

As mentioned above, the core can also generate burst accesses. In the case of non-blocking burst accesses, the *id* and *ready* parameters will be arrays instead of single values.

The description above covers basic operation of AT_AHB_MST. Please refer to the *grlib.at_ahb_mst_pkg* package located at *lib/grlib/atf/at_ahb_mst_pkg.vhd* to see all available procedure calls. Each call and its parameters are documented in the package.

7.3 AT AHB Slave

7.3.1 Description

The AT AHB Slave (AT_AHB_SLV) is a non-synthesizable AHB slave core with a debug interface that allows insertion of custom AHB replies and access to the core's internal memory structures.

7.3.2 Initialization and Instantiation

The component for the slave is defined in the package *grlib.at_pkg* and the procedure calls used to access the slave via its debug interface are available in the package *grlib.at_ahb_slv_pkg*. In order to instantiate the slave, the following libraries should be included:

```

library ieee;
use ieee.std_logic_1164.all;
library grlib;
use grlib.amba.all;
use grlib.at_pkg.all;
use grlib.at_ahb_slv_pkg.all;

```

The component for AT_AHB_SLV has the following interface:

```

component at_ahb_slv is
  generic (
    hindex      : integer := 0;      -- Slave index
    bank0addr   : integer := 0;
    bank0mask   : integer := 0;
    bank0type   : integer := 0;      -- 0: memory area 1: I/O area
    bank0cache  : integer := 0;      -- Cachable
    bank0prefetch : integer := 0;    -- Prefetchable
    bank0ws     : integer := 0;      -- Waitstates
    bank0rws    : integer := 0;      -- Random wait states 'ws' is the maximum
    bank0dataload : integer := 0;    -- Load data from file
    bank0datafile : string := "none"; -- Initial data for bank
    bankladdr   : integer := 0;
    banklmask   : integer := 0;
    bankltype   : integer := 0;      -- 0: memory area 1: I/O area
  );
end component;

```

LEON/GRLIB Guide

```

bank1cache      : integer := 0;      -- Cachable
bank1prefetch   : integer := 0;      -- Prefetchable
bank1lws        : integer := 0;      -- Waitstates
bank1rws        : integer := 0;      -- Random wait states 'ws' is the maximum
bank1dataload   : integer := 0;      -- Load data from file
bank1datafile   : string := "none";  -- Initial data for bank
bank2addr       : integer := 0;
bank2mask       : integer := 0;
bank2type       : integer := 0;      -- 0: memory area 1: I/O area
bank2cache      : integer := 0;      -- Cachable
bank2prefetch   : integer := 0;      -- Prefetchable
bank2lws        : integer := 0;      -- Waitstates
bank2rws        : integer := 0;      -- Random wait states 'ws' is the maximum
bank2dataload   : integer := 0;      -- Load data from file
bank2datafile   : string := "none";  -- Initial data for bank
bank3addr       : integer := 0;
bank3mask       : integer := 0;
bank3type       : integer := 0;      -- 0: memory area 1: I/O area
bank3cache      : integer := 0;      -- Cachable
bank3prefetch   : integer := 0;      -- Prefetchable
bank3lws        : integer := 0;      -- Waitstates
bank3rws        : integer := 0;      -- Random wait states 'ws' is the maximum
bank3dataload   : integer := 0;      -- Load data from file
bank3datafile   : string := "none";  -- Initial data for bank
grrlibdatamux   : integer := 1      -- GRLIB AMBA data MUX:ing
);
port (
  rstn : in  std_ulogic;
  clk  : in  std_ulogic;
  ahbsi : in  ahb_slv_in_type;
  ahbso : out ahb_slv_out_type;
  dbg_i : in  at_slv_dbg_in_type;
  dbg_o : out at_slv_dbg_out_type
);
end component;
```

The *hindex* generic must match the bus index in the same way as for other GRLIB cores. The *grrlib-datamux* generic decides if the core should use AMBA compliant data multiplexing (*grrlibdatamux* => 0) or the simplified data multiplexing scheme (*grrlibdatamux* => 1) used in GRLIB (see the GRLIB IP Library User's Manual, *grrlib.pdf*, for details).

For use in a normal GRLIB system, the default value is recommended. The other generics define the size and behavior of the, up to, four available AHB memory areas (banks). Each bank is configured via a set of generics described in the table below:

VHDL generic	Description
bank*addr	Bank base address. Set in the same manner as for all GRLIB AHB slaves
bank*mask	Bank mask. Decides how many of the bank*addr bits that are matched against the incoming AMBA HADDR and thereby also determines the size of the memory area.
bank*type	Selects if the bank is an AHB memory area or an AHB I/O area. The AT_AHB_SLV package defines to constants that can be used to select the type: AT_AHBSLV_MEM and AT_AHBSLV_IO.
bank*cache	Determines if bank is cacheable. This value is only used when banktype is set to AT_AHBSLV_MEM.
bank*prefetch	Determines if the bank is prefetchable. This value is only used when banktype is set to AT_AHBSLV_MEM.
bank*ws	Number of wait states that the core will insert on each access to the bank.
bank*rws	Enables random wait states. If this generic is set to AT_AHBSLV_RANDOM_WS, the core will insert between 0 and bank*ws wait states on each access. If this generic is set to AT_AHBSLV_FIXED_WS the core will always insert bank*ws wait states.
bank*dataload	If this generics is non-zero, the core will load initial memory data from the SREC file specified by bank*datafile.
bank*datafile	See above.

TABLE 11. AT_AHB_SLV VHDL generics

LEON/GRLIB Guide

An example instantiation of AT_AHB_SLV can be found in *verification/at/at_tb.vhd*. At the top of the file the libraries mentioned above is included. The signals used to make accesses to AT_AHB_SLV's debug interface are created with:

```
signal dbgi : at_slv_dbg_in_type;
signal dbgo : at_slv_dbg_out_type;
```

An example instantiation of AT_AHB_SLV looks like:

```
ahbslv0 : at_ahb_slv
  generic map (
    hindex => 0,
    -- Bank 0 configuration;
    bank0addr => 16#000#,
    bank0mask => 16#FFF#,
    bank0type => AT_AHBSLV_MEM,
    bank0cache => 1,
    bank0prefetch => 1,
    bank0ws => 1,
    bank0rws => AT_AHBSLV_FIXED_WS,
    bank0dataload => 0,
    bank0datafile => "none")
  port map (
    rstn => rstn, clk => clk,
    ahbsi => ahbsi, ahbso => ahbso(0),
    dbgi => dbgi, dbgo => dbgo);
```

After the *rstn* signal has gone high the core will be ready to handle incoming AMBA accesses. If no file is used to initialize the memory, all memory position will contain 'U'.

7.3.3 Controlling AT_AHB_SLV

When the slave has left system reset (*rstn* input is high), the procedures defined in *grlib.at_ahb_slv_pkg* (*lib/grlib/atf/at_ahb_slv_pkg.vhd*) can be used to control the slave's behavior and to access the slave's internal memory.

Accesses to the slave's internal memory are made via the *ahbslv_read(..)* and *ahbslv_write(..)* procedures. These procedures have the following interface:

```
-- Subprogram: ahbslv_write
-- Description: Write data to slave memory. The input address is masked and
--              only the valid bits are used. This means that the full AMBA
--              address can be used and the caller does not have to subtract
--              the bank start address.
procedure ahbslv_write (
  constant address : in  std_logic_vector(ADDR_R);
  constant data    : in  std_logic_vector;
  constant bank    : in  integer;
  signal  dbgi     : out at_slv_dbg_in_type;
  signal  dbgo     : in  at_slv_dbg_out_type);

-- Subprogram: ahbslv_read
-- Description: Read data from slave memory. The input address is masked and
--              only the valid bits are used. This means that the full AMBA
--              address can be used and the caller does not have to subtract
--              the bank start address.
procedure ahbslv_read (
  constant address : in  std_logic_vector(ADDR_R);
  variable data    : out std_logic_vector;
  constant bank    : in  integer;
  signal  dbgi     : out at_slv_dbg_in_type;
  signal  dbgo     : in  at_slv_dbg_out_type);
```

These functions are useful quickly initializing memory or to check the result of AMBA accesses made to the slave without generating traffic on the AMBA AHB bus. The width of the vector assigned to the data parameter determines the size of the access. The width of the address vector input must be 32 bits (31 downto 0).

A common use of AT_AHB_SLV is to specify special responses in order to test the behavior of AHB masters in the system. Custom responses can be inserted with the *ahbslv_response(..)* procedure. This procedure name is overloaded and variants with a different number of parameters exist. The most versatile *ahbslv_response(..)* procedure is:

```
procedure ahbslv_response (
  constant address_start : in  std_logic_vector(ADDR_R);
  constant address_stop  : in  std_logic_vector(ADDR_R);
  constant bank          : in  integer;
  constant response      : in  std_logic_vector(1 downto 0);
  constant data          : in  std_logic_vector;
  constant master        : in  integer range 0 to NAHBMST-1;
```

LEON/GRLIB Guide

```

constant anymst      : in  boolean;
variable id          : out integer;
signal  dbg_i        : out at_slv_dbg_in_type;
signal  dbg_o        : in  at_slv_dbg_out_type;
constant ws          : in  integer := 0;
constant repeat      : in  integer := 1;
constant count       : in  integer := 1;
constant splitcnt    : in  integer := 5;
constant mem_access  : in  boolean := false;
constant read_response : in  boolean := true;
constant write_response : in  boolean := true;
constant lock        : in  boolean := false;
constant delay       : in  integer := 0;
constant hprot       : in  std_logic_vector(3 downto 0);
constant anyhprot    : in  boolean;

```

The parameters are documented in the *grlib.at_ahb_slv_pkg* package. Note that several parameters have default values, this means that they do not have to be assigned when using the procedure. A selection of available AT_AHB_SLV procedures are listed in table 12. All procedures are further documented in the *grlib.at_ahb_slv_pkg* package located at *lib/grlib/atf/at_ahb_slv_pkg.vhd*.

Procedure name	Description
ahbslv_response	Inserts a customized response into the slaves response queue. If two responses are inserted for the same address (range), the first response to be inserted will be the first given. Several overloaded versions exist giving the ability to, for instance, only replying to accesses from a specific master that have a specific HPROT value. When a response is inserted, a unique identifier for that response is returned.
ahbslv_response_status	Used to determine if a response with a specified identifier is in the slave's response queue.
ahbslv_response_remove	Removes a response with a specified identifier from the slave's response queue.
ahbslv_response_clear	Removes all queue responses in the slave or only for a specified bank.
ahbslv_response_unlock	A response inserted with ahbslv_response(..) can be "locked" which means that it will be valid for an unlimited number of accesses. This procedure can be used to "unlock" the response, removing it from the slave.
ahbslv_waitforaccess	This procedure will block until an access has been made to a specified memory address.
ahbslv_waitforcomplete	This procedure will block until a queued response has been triggered and removed from the slave's response queue.
ahbslv_setconfig	Changes the default behavior of AHB slave model. Can be used to configure wait states, random wait states, random RETRY and SPLIT responses, etc.
ahbslv_getconfig	Reads the current default behavior of the slave.
ahbslv_enable_split	Enables SPLIT responses with a specified probability.
ahbslv_disable_split	Disables SPLIT responses.
ahbslv_enable_retry	Enables RETRY responses with a specified probability.
ahbslv_disable_retry	Disables RETRY responses.
ahbslv_set_ws	Sets the default number of wait states to be inserted by the slave.
ahbslv_get_ws	Gets the default number of wait states inserted by the slave.

TABLE 12. Selection of AT_AHB_SLV procedures

7.4 AT AHB Controller

7.4.1 Description

The AT AHB Controller (AT_AHB_CTRL) is a non-synthesizable AHB arbiter/controller. Compared to the standard GRLIB AHBCTRL core, AT_AHB_CTRL supports early burst termination and forced re-arbitration.

7.4.2 Usage

In order to instantiate the controller, the following libraries should be included:

```
library ieee;
use ieee.std_logic_1164.all;
library grib;
use grib.amba.all;
use grib.at_pkg.all;
```

The component for AT_AHB_CTRL has the following interface:

```
component at_ahb_ctrl is
  generic (
    defmast      : integer := 0;           -- default master
    split        : integer := 0;           -- split support
    rrobin       : integer := 0;           -- round-robin arbitration
    timeout      : integer range 0 to 255 := 0; -- HREADY timeout
    ioaddr       : ahb_addr_type := 16#fff#; -- I/O area MSB address
    iomask       : ahb_addr_type := 16#fff#; -- I/O area address mask
    cfgaddr      : ahb_addr_type := 16#ff0#; -- config area MSB address
    cfgmask      : ahb_addr_type := 16#ff0#; -- config area address mask
    nahbm        : integer range 1 to NAHBMST := NAHBMST; -- number of masters
    nahbs        : integer range 1 to NAHBSLV := NAHBSLV; -- number of slaves
    ioen         : integer range 0 to 15 := 1; -- enable I/O area
    disirq       : integer range 0 to 1 := 0; -- disable interrupt routing
    fixbrst      : integer range 0 to 1 := 0; -- support fix-length bursts
    debug        : integer range 0 to 2 := 2; -- report cores to console
    fppn         : integer range 0 to 1 := 0; -- full PnP configuration decoding
    icheck       : integer range 0 to 1 := 1;
    devid        : integer := 0;           -- unique device ID
    enbusmon     : integer range 0 to 1 := 0; --enable bus monitor
    assertwarn   : integer range 0 to 1 := 0; --enable assertions for warnings
    asserterr    : integer range 0 to 1 := 0; --enable assertions for errors
    hmstdisable  : integer := 0; --disable master checks
    hslvdisable  : integer := 0; --disable slave checks
    arbdisable   : integer := 0; --disable arbiter checks
    mprio        : integer := 0; --master with highest priority
    mcheck       : integer := 1; --check memory map for intersects
    enebterm     : integer := 0; --enable early burst termination
    ebprob       : integer := 10; --probability setting for of early bursttermination
    ccheck       : integer range 0 to 1 := 1; --perform sanity checks on pnp config
    acdm         : integer := 0; --AMBA compliant data muxing (for hsize > word)
  );
  port (
    rst      : in  std_ulogic;
    clk      : in  std_ulogic;
    msti     : out ahb_mst_in_type;
    msto     : in  ahb_mst_out_vector;
    slvi     : out ahb_slv_in_type;
    slvo     : in  ahb_slv_out_vector;
    testen   : in  std_ulogic := '0';
    testrst  : in  std_ulogic := '1';
    scanen   : in  std_ulogic := '0';
    testoen  : in  std_ulogic := '1';
    doarb    : in  std_ulogic := '0'
  );
end component;
```

Most of the core's VHDL generics are the same as for the AHBCTRL core. Two generics have been added: *enebterm* and *ebprob*. When *enebterm* is set to a non-zero value the core may automatically terminate burst accesses early. The normal GRLIB arbiter, AHBCTRL, does not interrupt a burst by removing grant from a master. With *enebterm* /= 0 and *ebprob* set to 10 the probability of a burst being interrupted by AT_AHB_CTRL is about 0.10 in each cycle.

Bursts may also be terminated early by assertion of the *doarb* input signal. When *doarb* is asserted, the AHB arbiter will perform arbitration.

Use of AT_AHB_CTRL is primarily recommended when a core will be used in non-GRLIB systems. The GRLIB arbiter will never interrupt a burst access and it is not a strict requirement that a core can handle terminated bursts for the core to function in GRLIB.

8 Support

Frontgrade Gaisler AB provides support via support@gaisler.com for customers with support contracts. Limited free support is also provided by Frontgrade Gaisler engineers on the GRLIB community found at <https://discourse.grib.com/community/>.

Frontgrade Gaisler AB
Kungsgatan 12
411 19 Göteborg
Sweden
www.frontgrade.com/gaisler
sales@gaisler.com
T: +46 31 7758650

Frontgrade Gaisler AB, reserves the right to make changes to any products and services described herein at any time without notice. Consult the company or an authorized sales representative to verify that the information in this document is current before using this product. The company does not assume any responsibility or liability arising out of the application or use of any product or service described herein, except as expressly agreed to in writing by the company; nor does the purchase, lease, or use of a product or service from the company convey a license under any patent rights, copyrights, trademark rights, or any other of the intellectual rights of the company or of third parties. All information is provided as is. There is no warranty that it is correct or suitable for any purpose, neither implicit nor explicit.

Copyright © 2025 Frontgrade Gaisler AB